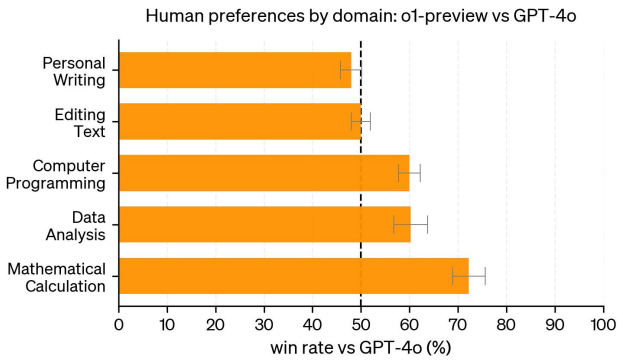

Figure 9: Preference Evaluation Results between GPT4o and GPT4-o1

very likely a significant overlap in workers who annotate the preferences as recently discussed [107]. Creative writing may also present an inherent "quality ceiling" for current LLM architectures, where all models encounter similar limitations in maintaining a consistent narrative voice and avoiding clichés—a pattern also observed by [112] in their analysis of narrative generation. Our syntactic pattern analysis (Section 5.3) reveals remarkably similar templates and phrasings across models, further supporting the hypothesis of shared training foundations. While our methodology focused on concrete, actionable edit categories, we acknowledge that more subtle stylistic differences or subjective "vibes" may exist between models not captured by our current taxonomy [31]. The consistency in writing quality across models ultimately points to broader questions about the current limitations of LLMs in creative writing tasks, rather than differences between specific implementations. Recent preference evaluation results from OpenAI's GPT4-o1 technical report [84] (See Figure 9) corroborate our findings where there are no significant differences in GPT4o vs GPT4o1 for personal writing or editing text.

8 Limitations

While our study provides valuable insights into improving LLM-generated text through expert editing, there are several limitations to consider. Our study was conducted with 18 MFA-trained creative writers. While this ensured a high level of expertise, it may limit the generalizability of our findings. Future research could expand the participant pool to encompass diverse cultural backgrounds and writing traditions. The editing data primarily comes from literary fiction and creative non-fiction, making the identified idiosyncrasies and editing strategies potentially less applicable to other genres like technical writing, journalism, or scientific writing. Future work can expand on the line of work by including a broader range of writing styles and purposes. The selected LLMs (GPT-4, Claude 3.5 Sonnet, and Llama 3.1) are among the most advanced models, but they might not fully represent the entire spectrum of AI writing abilities. It should be noted that the evaluation of writing quality is inherently subjective, even with multiple annotators and inter-annotator agreement calculations. Experts may disagree on what constitutes an improvement, potentially influencing our results and their interpretation.Our automated methods for detecting and rewriting problematic spans relied on few-shot learning with a limited number of examples. While this approach showed promise, it may not fully capture the complexity and nuance of expert editing and training a model on the entire LAMP Corpus or additional data is required.

It should be noted that while paragraph-level editing provides a balance between granularity and context, it may miss broader structural or thematic issues that become apparent only when considering longer pieces of writing. Last but not least, our study relied on professional writers editing AI-generated text for monetary compensation, which may have influenced the quality and nature of the edits. Editing one's own work typically involves more personal investment than editing text for pay, potentially leading to less motivation for substantial improvements [13]. Additionally, the repetitive nature of editing multiple AI-generated paragraphs could lead to fatigue, especially if the content is perceived as uninteresting or lacking in creativity. This fatigue could result in less thorough or thoughtful edits as the task progresses. Finally, another potential limitation of LLM-based editing is the risk of hallucinated and factually inconsistent information being introduced during the editing process. We did not study hallucinations as part of our work, due to the chosen domain focus of fictional writing which has less stringent factuality requirements, yet prior work has documented that LLM-based text editing can introduce factual errors [61].

9 Conclusion

In this work, we present a comprehensive approach to mitigating idiosyncrasies and improving human-AI alignment in the writing process through expert editing. We i) develop a taxonomy of edit categories grounded in established writing practices, ii) create the LAMP corpus containing over 8,000 fine-grained edits by professional writers on LLM-generated text, and iii) design methods for automatic detection and rewriting of problematic spans. Our analysis reveals several key findings. Professional writers identify consistent categories of edits needed to improve AI writing. Surprisingly, there are no significant differences in perceived writing quality or types of edits needed across texts generated by different large language models (GPT-4, Claude 3.5, Llama 3.1). Automated methods using few-shot prompting can detect and rewrite problematic spans in LLM-generated text, though far from matching human expert performance. Finally, in terms of preference evaluations, writers consistently rank text edited by other writers highest, followed by LLM-edited text, with unedited LLM-generated text ranking lowest. As AI text generation becomes more prevalent, developing robust editing and alignment techniques will be crucial to ensure AI systems produce high-quality writing that meets human standards and enhances creativity and linguistic diversity.

References

[1] 2018. Purple prose. https://en.wikipedia.org/wiki/Purple_prose
[2] Teresa M Amabile. 1982. Social psychology of creativity: A consensual assessment technique. *Journal of personality and social psychology* 43, 5 (1982), 997.

[3] Barrett R Anderson, Jash Hemant Shah, and Max Kreminski. 2024. Homogenization Effects of Large Language Models on Human Creative Ideation. In *Proceedings of the 16th Conference on Creativity & Cognition* (Chicago, IL, USA) *(C&C '24)*. Association for Computing Machinery, New York, NY, USA, 413–425. https://doi.org/10.1145/3635636.3656204

[4] Barrett R Anderson, Jash Hemant Shah, and Max Kreminski. 2024. Homogenization effects of large language models on human creative ideation. In *Proceedings of the 16th Conference on Creativity & Cognition*. 413–425.

[5] Anthropic. 2024. Claude-3.5-Sonnet. (2024). https://www.anthropic.com/news/claude-3-5-sonnet

[6] Kenneth C. Arnold, Krysta Chauncey, and Krzysztof Z. Gajos. 2020. Predictive text encourages predictable writing. In *Proceedings of the 25th International Conference on Intelligent User Interfaces* (Cagliari, Italy) *(IUI '20)*. Association for Computing Machinery, New York, NY, USA, 128–138. https://doi.org/10.1145/3377325.3377523

[7] Michael S. Bernstein, Greg Little, Robert C. Miller, Björn Hartmann, Mark S. Ackerman, David R. Karger, David Crowell, and Katrina Panovich. 2015. Soylent: a word processor with a crowd inside. *Commun. ACM* 58, 8 (jul 2015), 85–94. https://doi.org/10.1145/2791285

[8] John R Bormuth. 1966. Readability: A new approach. *Reading research quarterly* (1966), 79–132.

[9] Robert J Bracewell et al. 1978. The development of audience awareness in writing. (1978).

[10] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. 2020. Language Models are Few-Shot Learners. In *Advances in Neural Information Processing Systems*, H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin (Eds.), Vol. 33. Curran Associates, Inc., 1877–1901. https://proceedings.neurips.cc/paper_files/paper/2020/file/1457c0d6bfcb4967418bfb8ac142f64a-Paper.pdf

[11] Janet Burroway, Elizabeth Stuckey-French, and Ned Stuckey-French. 2019. *Writing fiction: A guide to narrative craft*. University of Chicago Press.

[12] Daniel Buschek, Martin Zürn, and Malin Eiband. 2021. The impact of multiple parallel phrase suggestions on email input and composition behaviour of native and non-native english writers. In *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*. 1–13.

[13] Kris Byron and Shalini Khazanchi. 2012. Rewards and creative performance: a meta-analytic test of theoretically derived hypotheses. *Psychological bulletin* 138, 4 (2012), 809.

[14] Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémy Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, et al. 2023. Open problems and fundamental limitations of reinforcement learning from human feedback. *arXiv preprint arXiv:2307.15217* (2023).

[15] Tuhin Chakrabarty, Philippe Laban, Divyansh Agarwal, Smaranda Muresan, and Chien-Sheng Wu. 2024. Art or Artifice? Large Language Models and the False Promise of Creativity. In *Proceedings of the CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) *(CHI '24)*. Association for Computing Machinery, New York, NY, USA, Article 30, 34 pages. https://doi.org/10.1145/3613904.3642731

[16] Tuhin Chakrabarty, Vishakh Padmakumar, Faeze Brahman, and Smaranda Muresan. 2024. Creativity Support in the Age of Large Language Models: An Empirical Study Involving Professional Writers. In *Proceedings of the 16th Conference on Creativity & Cognition* (Chicago, IL, USA) *(C &C '24)*. Association for Computing Machinery, New York, NY, USA, 132–155. https://doi.org/10.1145/3635636.3656201

[17] Anton Chekhov. 1984. *The selected letters of Anton Chekhov*. Macmillan.

[18] Eunsol Choi, Jennimaria Palomaki, Matthew Lamm, Tom Kwiatkowski, Dipanjan Das, and Michael Collins. 2021. Decontextualization: Making Sentences Stand-Alone. *Transactions of the Association for Computational Linguistics* 9 (2021), 447–461. https://doi.org/10.1162/tacl_a_00377

[19] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. 2017. Deep reinforcement learning from human preferences. *Advances in neural information processing systems* 30 (2017).

[20] Gregory Currie. 1990. *The nature of fiction*. Cambridge University Press.

[21] Hai Dang, Karim Benharrak, Florian Lehmann, and Daniel Buschek. 2022. Beyond Text Generation: Supporting Writers with Continuous Automatic Text Summaries. In *Proceedings of the 35th Annual ACM Symposium on User Interface Software and Technology* (Bend, OR, USA) *(UIST '22)*. Association for Computing Machinery, New York, NY, USA, Article 98, 13 pages. https://doi.org/10.1145/3526113.3545672

[22] Johannes Daxenberger and Iryna Gurevych. 2013. Automatically Classifying Edit Categories in Wikipedia Revisions. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, David Yarowsky, Timothy Baldwin, Anna Korhonen, Karen Livescu, and Steven Bethard (Eds.). Association for Computational Linguistics, Seattle, Washington, USA, 578–589. https://aclanthology.org/D13-1055

[23] Marshall L Dermer, Shannon L Lopez, and Paul A Messling. 2009. Fluency training a writing skill: Editing for concision. *The Psychological Record* 59 (2009), 3–20.

[24] Paramveer S. Dhillon, Somayeh Molaei, Jiaqi Li, Maximilian Golub, Shaochun Zheng, and Lionel Peter Robert. 2024. Shaping Human-AI Collaboration: Varied Scaffolding Levels in Co-writing with Language Models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) *(CHI '24)*. Association for Computing Machinery, New York, NY, USA, Article 1044, 18 pages. https://doi.org/10.1145/3613904.3642134

[25] Karel D'Oosterlinck, Winnie Xu, Chris Develder, Thomas Demeester, Amanpreet Singh, Christopher Potts, Douwe Kiela, and Shikib Mehri. 2024. Anchored Preference Optimization and Contrastive Revisions: Addressing Underspecification in Alignment. *arXiv preprint arXiv:2408.06266* (2024).

[26] Mark Doty. 2014. *The art of description: World into word*. Graywolf Press.

[27] Yao Dou, Philippe Laban, Claire Gardent, and Wei Xu. 2023. Automatic and Human-AI Interactive Text Generation. *arXiv preprint arXiv:2310.03878* (2023).

[28] Yao Dou, Philippe Laban, Claire Gardent, and Wei Xu. 2024. Automatic and Human-AI Interactive Text Generation (with a focus on Text Simplification and Revision). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 5: Tutorial Abstracts)*, Luis Chiruzzo, Hung-yi Lee, and Leonardo Ribeiro (Eds.). Association for Computational Linguistics, Bangkok, Thailand, 3–4. https://doi.org/10.18653/v1/2024.acl-tutorials.2

[29] Wanyu Du, Vipul Raheja, Dhruv Kumar, Zae Myung Kim, Melissa Lopez, and Dongyeop Kang. 2022. Understanding Iterative Revision from Human-Written Text. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (Eds.). Association for Computational Linguistics, Dublin, Ireland, 3573–3590. https://doi.org/10.18653/v1/2022.acl-long.250

[30] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783* (2024).

[31] Lisa Dunlap, Krishna Mandal, Trevor Darrell, Jacob Steinhardt, and Joseph E Gonzalez. 2024. VibeCheck: Discover and Quantify Qualitative Differences in Large Language Models. *arXiv preprint arXiv:2410.12851* (2024).

[32] Jane Dwivedi-Yu, Timo Schick, Zhengbao Jiang, Maria Lomeli, Patrick Lewis, Gautier Izacard, Edouard Grave, Sebastian Riedel, and Fabio Petroni. 2022. Editeval: An instruction-based benchmark for text improvements. *arXiv preprint arXiv:2209.13331* (2022).

[33] Lester Faigley and Stephen Witte. 1981. Analyzing revision. *College Composition & Communication* 32, 4 (1981), 400–414.

[34] Felix Faltings, Michel Galley, Gerold Hintz, Chris Brockett, Chris Quirk, Jianfeng Gao, and Bill Dolan. 2021. Text Editing by Command. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, Kristina Toutanova, Anna Rumshisky, Luke Zettlemoyer, Dilek Hakkani-Tur, Iz Beltagy, Steven Bethard, Ryan Cotterell, Tanmoy Chakraborty, and Yichao Zhou (Eds.). Association for Computational Linguistics, Online, 5259–5274. https://doi.org/10.18653/v1/2021.naacl-main.414

[35] Andy P Field. 2005. Kendall's coefficient of concordance. *Encyclopedia of statistics in behavioral science* 2 (2005), 1010–11.

[36] David Fishelov. 1990. Types of character, characteristics of types. *Style* (1990), 422–439.

[37] Kristie S Fleckenstein. 1991. Inner Sight: Imagery and Emotion in Writing Engagement. *Teaching English in the Two-Year College* 18, 3 (1991), 210–16.

[38] Nigel Fountain. 2012. *Clichés: Avoid them like the plague*. Michael O'Mara Books.

[39] Iason Gabriel, Arianna Manzini, Geoff Keeling, Lisa Anne Hendricks, Verena Rieser, Hasan Iqbal, Nenad Tomašev, Ira Ktena, Zachary Kenton, Mikel Rodriguez, et al. 2024. The ethics of advanced ai assistants. *arXiv preprint arXiv:2404.16244* (2024).

[40] Ge Gao, Alexey Taymanov, Eduardo Salinas, Paul Mineiro, and Dipendra Misra. 2024. Aligning llm agents by learning latent preference from user edits. *arXiv preprint arXiv:2404.15269* (2024).

[41] Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, et al. 2020. The pile: An 800gb dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027* (2020).

[42] Yang Gao, Dana Alon, and Donald Metzler. 2024. Impact of preference noise on the alignment performance of generative language models. *arXiv preprint arXiv:2404.09824* (2024).

[43] Katy Ilonka Gero, Vivian Liu, and Lydia Chilton. 2022. Sparks: Inspiration for science writing using language models. In *Proceedings of the 2022 ACM Designing Interactive Systems Conference*. 1002–1019.

[44] Sayan Ghosh, Tejas Srinivasan, and Swabha Swayamdipta. 2024. Compare without Despair: Reliable Preference Evaluation with Generation Separability. *arXiv preprint arXiv:2407.01878* (2024).

[45] Yvette Graham, Timothy Baldwin, Alistair Moffat, and Justin Zobel. 2013. Continuous measurement scales in human evaluation of machine translation. In *Proceedings of the 7th Linguistic Annotation Workshop and Interoperability with Discourse*. 33–41.

[46] Harry Guinness. 2024. How to Edit Your Own Writing. (2024). https://www.nytimes.com/2020/04/07/smarter-living/how-to-edit-your-own-writing.html

[47] John R Hayes, Linda Flower, Karen A Schriver, James Stratman, Linda Carey, et al. 1987. Cognitive processes in revision. *Advances in applied psycholinguistics* 2 (1987), 176–240.

[48] Laraine Herring. 2007. *Writing Begins with the Breath: Embodying Authentic Voice*. Shambhala Publications.

[49] Jess Hohenstein and Malte Jung. 2020. AI as a moral crumple zone: The effects of AI-mediated communication on attribution and trust. *Computers in Human Behavior* 106 (2020), 106190.

[50] Daphne Ippolito, Liam Dugan, Emily Reif, Ann Yuan, Andy Coenen, and Chris Callison-Burch. 2022. The Case for a Single Model that can Both Generate Continuations and Fill-in-the-Blank. In *Findings of the Association for Computational Linguistics: NAACL 2022*, Marine Carpuat, Marie-Catherine de Marneffe, and Ivan Vladimir Meza Ruiz (Eds.). Association for Computational Linguistics, Seattle, United States, 2421–2432. https://doi.org/10.18653/v1/2022.findings-naacl.185

[51] Daphne Ippolito, Ann Yuan, Andy Coenen, and Sehmon Burnam. 2022. Creative Writing with an AI-Powered Writing Assistant: Perspectives from Professional Writers. *arXiv preprint arXiv:2211.05030* (2022).

[52] Maurice Jakesch, Advait Bhat, Daniel Buschek, Lior Zalmanson, and Mor Naaman. 2023. Co-Writing with Opinionated Language Models Affects Users' Views. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems* (Hamburg, Germany) *(CHI '23)*. Association for Computing Machinery, New York, NY, USA, Article 111, 15 pages. https://doi.org/10.1145/3544548.3581196

[53] Chao Jiang, Wei Xu, and Samuel Stevens. 2022. arXivEdits: Understanding the Human Revision Process in Scientific Writing. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. 9420–9435.

[54] Makena Kelly. 2023. *Books3 is the Internet's Most Controversial AI Dataset*. https://www.theatlantic.com/technology/archive/2023/09/books3-database-generative-ai-training-copyright-infringement/675363/ Accessed: 2024-11-18.

[55] Zae Myung Kim, Wanyu Du, Vipul Raheja, Dhruv Kumar, and Dongyeop Kang. 2022. Improving Iterative Text Revision by Learning Where to Edit from Other Revision Tasks. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 9986–9999. https://doi.org/10.18653/v1/2022.emnlp-main.678

[56] Stephen King. 2000. *On writing: A memoir of the craft*. Simon and Schuster.

[57] George R Klare. 1974. Assessing readability. *Reading research quarterly* (1974), 62–102.

[58] Jon Kleinberg and Manish Raghavan. 2021. Algorithmic monoculture and social welfare. *Proceedings of the National Academy of Sciences* 118, 22 (2021), e2018340118.

[59] Dmitry Kobak, Rita González Márquez, Emőke-Ágnes Horvát, and Jan Lause. 2024. Delving into ChatGPT usage in academic writing through excess vocabulary. *arXiv preprint arXiv:2406.07016* (2024).

[60] Anežka Kuzmičová. 2014. Literary narrative and mental imagery: A view from embodied cognition. *Style* 48, 3 (2014), 275–293.

[61] Philippe Laban, Jesse Vig, Marti A Hearst, Caiming Xiong, and Chien-Sheng Wu. 2023. Beyond the chat: Executable and verifiable text-editing with llms. *arXiv preprint arXiv:2309.15337* (2023).

[62] Philippe Laban, Jesse Vig, Wojciech Kryscinski, Shafiq Joty, Caiming Xiong, and Chien-Sheng Wu. 2023. SWiPE: A Dataset for Document-Level Simplification of Wikipedia Pages. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 10674–10695. https://doi.org/10.18653/v1/2023.acl-long.596

[63] Janet Lane and Ellen Lange. 1999. *Writing clearly: An editing guide*. Heinle & Heinle Boston, MA.

[64] Mina Lee, Katy Ilonka Gero, John Joon Young Chung, Simon Buckingham Shum, Vipul Raheja, Hua Shen, Subhashini Venugopalan, Thiemo Wambsganss, David Zhou, Emad A Alghamdi, et al. 2024. A Design Space for Intelligent and Interactive Writing Assistants. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*. 1–35.

[65] Mina Lee, Percy Liang, and Qian Yang. 2022. Coauthor: Designing a human-ai collaborative writing dataset for exploring language model capabilities. In *Proceedings of the 2022 CHI conference on human factors in computing systems*. 1–19.

[66] Vladimir I. Levenshtein. 1966. Binary Codes Capable of Correcting Deletions, Insertions and Reversals. *Soviet physics. Doklady* 10 (1966), 707–710.

[67] Xian Li, Ping Yu, Chunting Zhou, Timo Schick, Luke Zettlemoyer, Omer Levy, Jason Weston, and Mike Lewis. 2023. Self-alignment with instruction backtranslation. *arXiv preprint arXiv:2308.06259* (2023).

[68] Zhuoyan Li, Chen Liang, Jing Peng, and Ming Yin. 2024. The Value, Benefits, and Concerns of Generative AI-Powered Assistance in Writing. In *Proceedings of the CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) *(CHI '24)*. Association for Computing Machinery, New York, NY, USA, Article 1048, 25 pages. https://doi.org/10.1145/3613904.3642625

[69] Chin-Yew Lin. 2004. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*. 74–81.

[70] Ruibo Liu, Chenyan Jia, Ge Zhang, Ziyu Zhuang, Tony Liu, and Soroush Vosoughi. 2022. Second thoughts are best: Learning to re-align with human values from text edits. *Advances in Neural Information Processing Systems* 35 (2022), 181–196.

[71] Susan Peck MacDonald. 1986. Specificity in context: some difficulties for the inexperienced writer. *College Composition & Communication* 37, 2 (1986), 178–194.

[72] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, et al. 2024. Self-refine: Iterative refinement with self-feedback. *Advances in Neural Information Processing Systems* 36 (2024).

[73] Mounica Maddela, Yao Dou, David Heineman, and Wei Xu. 2022. LENS: A learnable evaluation metric for text simplification. *arXiv preprint arXiv:2212.09739* (2022).

[74] Guillermo Marco, Julio Gonzalo, Ramón del Castillo, and María Teresa Mateo Girona. 2024. Pron vs Prompt: Can Large Language Models already Challenge a World-Class Fiction Author at Creative Text Writing? *arXiv preprint arXiv:2407.01119* (2024).

[75] John McPhee. 2015. Omission. *The New Yorker* (14 September 2015). https://www.newyorker.com/magazine/2015/09/14/omission Accessed: [Insert access date here].

[76] Bonnie JF Meyer. 2003. Text coherence and readability. *Topics in language Disorders* 23, 3 (2003), 204–224.

[77] Piotr Mirowski, Juliette Love, Kory Mathewson, and Shakir Mohamed. 2024. A Robot Walks into a Bar: Can Language Models Serve as Creativity SupportTools for Comedy? An Evaluation of LLMs' Humour Alignment with Comedians. In *The 2024 ACM Conference on Fairness, Accountability, and Transparency*. 1622–1636.

[78] Piotr Mirowski, Kory W. Mathewson, Jaylen Pittman, and Richard Evans. 2023. Co-Writing Screenplays and Theatre Scripts with Language Models: Evaluation by Industry Professionals. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems* (Hamburg, Germany) *(CHI '23)*. Association for Computing Machinery, New York, NY, USA, Article 355, 34 pages. https://doi.org/10.1145/3544548.3581225

[79] Marc Moens. 1987. Tense, aspect and temporal reference. (1987).

[80] Michael Nebeling, Alexandra To, Anhong Guo, Adrian A. de Freitas, Jaime Teevan, Steven P. Dow, and Jeffrey P. Bigham. 2016. WearWrite: Crowd-Assisted Writing from Smartwatches. In *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems* (San Jose, California, USA) *(CHI '16)*. Association for Computing Machinery, New York, NY, USA, 3834–3846. https://doi.org/10.1145/2858036.2858169

[81] Nancy Nelson and Montserrat Castelló. 2012. Academic writing and authorial voice. In *University writing: Selves and texts in academic societies*. Brill, 33–51.

[82] William Noble. 2013. *Show Don't Tell: A Writer's Guide*. Vol. 6. The Write Thought, Inc.

[83] OpenAI. 2024. GPT-4o. (2024). https://openai.com/index/hello-gpt-4o/

[84] OpenAI. 2024. *Learning to Reason with LLMs*. Technical Report. OpenAI. https://openai.com/index/learning-to-reason-with-llms/

[85] George Orwell. 2013. *Politics and the English language*. Penguin UK.

[86] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. 2022. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems* 35 (2022), 27730–27744.

[87] Vishakh Padmakumar and He He. 2023. Does Writing with Language Models Reduce Content Diversity? *arXiv preprint arXiv:2309.05196* (2023).

[88] Jane Pan, He He, Samuel R Bowman, and Shi Feng. 2024. Spontaneous Reward Hacking in Iterative Self-Refinement. *arXiv preprint arXiv:2407.04549* (2024).

[89] Liangming Pan, Michael Saxon, Wenda Xu, Deepak Nathani, Xinyi Wang, and William Yang Wang. 2024. Automatically correcting large language models: Surveying the landscape of diverse automated correction strategies. *Transactions of the Association for Computational Linguistics* 12 (2024), 484–506.

[90] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*. 311–318.

[91] So Yeon Park and Sang Won Lee. 2023. Why "why"? The Importance of Communicating Rationales for Edits in Collaborative Writing. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems* (Hamburg, Germany) *(CHI '23)*. Association for Computing Machinery, New York, NY, USA, Article 616, 25 pages. https://doi.org/10.1145/3544548.3581345

[92] Roy D Pea and D Midian Kurland. 1987. Chapter 7: Cognitive Technologies for Writing. *Review of research in education* 14, 1 (1987), 277–326.