[93] Reid Pryzant, Richard Diehl Martinez, Nathan Dass, Sadao Kurohashi, Dan Jurafsky, and Diyi Yang. 2020. Automatically neutralizing subjective bias in text. In *Proceedings of the aaai conference on artificial intelligence*, Vol. 34. 480–489.

[94] Luqman M Rababah. 2022. Contextualization to enhance students' writing ability. *Theory and Practice in Language Studies* 12, 11 (2022), 2316–2321.

[95] Vipul Raheja, Dhruv Kumar, Ryan Koo, and Dongyeop Kang. 2023. CoEdIT: Text Editing by Task-Specific Instruction Tuning. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, Houda Bouamor, Juan Pino, and Kalika Bali (Eds.). Association for Computational Linguistics, Singapore, 5274–5291. https://doi.org/10.18653/v1/2023.findings-emnlp.350

[96] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. 2016. Squad: 100,000+ questions for machine comprehension of text. *arXiv preprint arXiv:1606.05250* (2016).

[97] Machel Reid and Graham Neubig. 2022. Learning to Model Editing Processes. In *Findings of the Association for Computational Linguistics: EMNLP 2022*, Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (Eds.). Association for Computational Linguistics, Abu Dhabi, United Arab Emirates, 3822–3832. https://doi.org/10.18653/v1/2022.findings-emnlp.280

[98] Mohi Reza, Nathan M Laundry, Ilya Musabirov, Peter Dushniku, Zhi Yuan "Michael" Yu, Kashish Mittal, Tovi Grossman, Michael Liut, Anastasia Kuzminykh, and Joseph Jay Williams. 2024. ABScribe: Rapid Exploration & Organization of Multiple Writing Variations in Human-AI Co-Writing Tasks using Large Language Models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) *(CHI '24)*. Association for Computing Machinery, New York, NY, USA, Article 1042, 18 pages. https://doi.org/10.1145/3613904.3641899

[99] Scott P. Robertson and John B. Black. 1983. Planning units in text editing behavior. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (Boston, Massachusetts, USA) *(CHI '83)*. Association for Computing Machinery, New York, NY, USA, 217–221. https://doi.org/10.1145/800045.801614

[100] Mary Beth Rosson. 1983. Patterns of experience in text editing. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (Boston, Massachusetts, USA) *(CHI '83)*. Association for Computing Machinery, New York, NY, USA, 171–175. https://doi.org/10.1145/800045.801604

[101] Keita Saito, Akifumi Wachi, Koki Wataoka, and Youhei Akimoto. 2023. Verbosity bias in preference labeling by large language models. *arXiv preprint arXiv:2310.10076* (2023).

[102] Marlene Scardamalia. 1983. The development of evaluative, diagnostic, and remedial capabilities in children's composing. *The psychology of written language: A developmental approach* (1983), 67–95.

[103] Chantal Shaib, Yanai Elazar, Junyi Jessy Li, and Byron C Wallace. 2024. Detection and Measurement of Syntactic Templates in Generated Text. *arXiv preprint arXiv:2407.00211* (2024).

[104] Lei Shu, Liangchen Luo, Jayakumar Hoskere, Yun Zhu, Yinxiao Liu, Simon Tong, Jindong Chen, and Lei Meng. 2024. Rewritelm: An instruction-tuned large language model for text rewriting. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 38. 18970–18980.

[105] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. 2024. AI models collapse when trained on recursively generated data. *Nature* 631, 8022 (2024), 755–759.

[106] Nancy Sommers. 1980. Revision strategies of student writers and experienced adult writers. *College Composition & Communication* 31, 4 (1980), 378–388.

[107] Wall Street Journal Staff. 2024. Alexandr Wang: Scale AI. *The Wall Street Journal* (2024). https://www.wsj.com/tech/ai/alexandr-wang-scale-ai-d7c6efd7 Accessed: 2024-11-18.

[108] Melanie Subbiah, Sean Zhang, Lydia B Chilton, and Kathleen McKeown. 2024. Reading Subtext: Evaluating Large Language Models on Short Story Summarization with Writers. *arXiv preprint arXiv:2403.01061* (2024).

[109] G Thomas Tanselle. 1976. The editorial problem of final authorial intention. *Studies in Bibliography* 29 (1976), 167–211.

[110] Jaime Teevan, Shamsi T. Iqbal, and Curtis von Veh. 2016. Supporting Collaborative Writing with Microtasks. In *Proceedings of the 2016 CHI Conference on Human Factors in Computing Systems* (San Jose, California, USA) *(CHI '16)*. Association for Computing Machinery, New York, NY, USA, 2657–2668. https://doi.org/10.1145/2858036.2858108

[111] David R Thomas. 2006. A general inductive approach for analyzing qualitative evaluation data. *American journal of evaluation* 27, 2 (2006), 237–246.

[112] Yufei Tian, Tenghao Huang, Miri Liu, Derek Jiang, Alexander Spangher, Muhao Chen, Jonathan May, and Nanyun Peng. 2024. Are Large Language Models Capable of Generating Human-Level Narratives? *arXiv preprint arXiv:2407.13248* (2024).

[113] Sherman W. Tyler, Steven Roth, and Timothy Post. 1982. The acquisition of text editing skills. In *Proceedings of the 1982 Conference on Human Factors in Computing Systems* (Gaithersburg, Maryland, USA) *(CHI '82)*. Association for Computing Machinery, New York, NY, USA, 324–325. https://doi.org/10.1145/800049.801803

[114] Manya Wadhwa, Xinyu Zhao, Junyi Jessy Li, and Greg Durrett. 2024. Learning to Refine with Fine-Grained Natural Language Feedback. In *Findings of the Association for Computational Linguistics: EMNLP 2024*. 12281–12308.

[115] Diyi Yang, Aaron Halfaker, Robert Kraut, and Eduard Hovy. 2017. Identifying Semantic Edit Intentions from Revisions in Wikipedia. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, Martha Palmer, Rebecca Hwa, and Sebastian Riedel (Eds.). Association for Computational Linguistics, Copenhagen, Denmark, 2000–2010. https://doi.org/10.18653/v1/D17-1213

[116] Zonghai Yao, Benjamin Schloss, and Sai Selvaraj. 2023. Improving Summarization with Human Edits. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*. 2604–2620.

[117] Ann Yuan, Andy Coenen, Emily Reif, and Daphne Ippolito. 2022. Wordcraft: story writing with large language models. In *Proceedings of the 27th International Conference on Intelligent User Interfaces*. 841–852.

[118] Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q Weinberger, and Yoav Artzi. 2019. Bertscore: Evaluating text generation with bert. *arXiv preprint arXiv:1904.09675* (2019).

[119] David Zhou and Sarah Sterman. 2024. Ai.llude: Investigating Rewriting AI-Generated Text to Support Creative Expression. In *Proceedings of the 16th Conference on Creativity & Cognition* (Chicago, IL, USA) *(C & C '24)*. Association for Computing Machinery, New York, NY, USA, 241–254. https://doi.org/10.1145/3635636.3656187

[120] Yilun Zhou, Caiming Xiong, Silvio Savarese, and Chien-Sheng Wu. 2024. Shared Imagination: LLMs Hallucinate Alike. *arXiv preprint arXiv:2407.16604* (2024).

[121] Daniel M Ziegler, Nisan Stiennon, Jeffrey Wu, Tom B Brown, Alec Radford, Dario Amodei, Paul Christiano, and Geoffrey Irving. 2019. Fine-tuning language models from human preferences. *arXiv preprint arXiv:1909.08593* (2019).

A Appendix

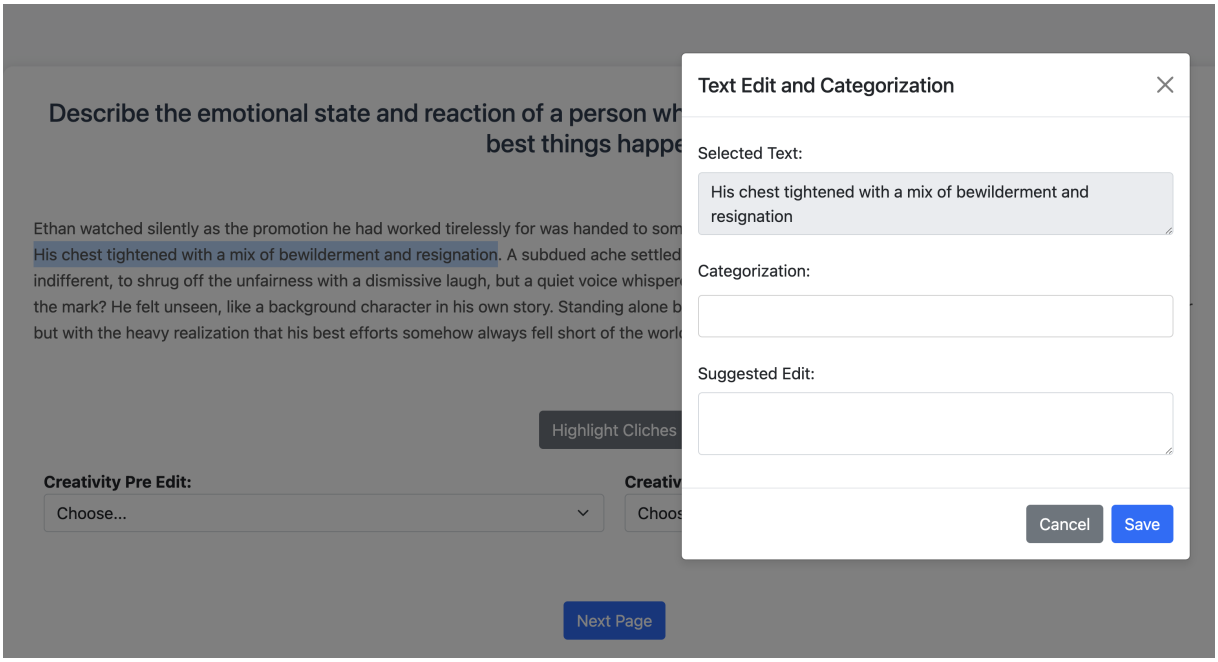

Figure 10: Interface for formative study to collect fine-grained labels

Instruction Prompt	Summarize this paragraph into a single sentence open-ended question.\n **{{paragraph}}**
	Summarize this paragraph into a single sentence open-ended instruction.\n **{{paragraph}}**
Response Prompt	Imagine you are a fiction writer for the NewYorker. Now write a paragraph (10-15 sentence) as a response to the following question. Try your best to be original, avoiding clichés or overused tropes. Do not use ornamental language and focus on nuance, simplicity, and subtext. Start directly with your response. \n **{{instruction}}**
	Imagine you are a writer for the New York Times Modern Love section. Now write a paragraph (10-15 sentence) as a response to the following question. Try your best to be original, avoiding clichés or overused tropes. Do not use ornamental language and focus on nuance, simplicity, and subtext.Start directly with your response \n **{{instruction}}**
	Imagine you are a writer for the New York Times Cooking section. Now write a paragraph (10-15 sentence) as a response to the following question. Try your best to be original, avoiding clichés or overused tropes. Do not use ornamental language and focus on nuance, simplicity, and subtext.Start directly with your response \n **{{instruction}}**
	Imagine you are a writer for the New York Times Travel section. Now write a paragraph (10-15 sentence) as a response to the following question. Try your best to be original, avoiding clichés or overused tropes. Do not use ornamental language and focus on nuance, simplicity, and subtext.Start directly with your response \n **{{instruction}}**
	Imagine you are a beloved female Internet advice columnist whose trademark is deeply felt and frank responses grounded in your own personal experience. Now write a paragraph (10-15 sentence) as a response to the following question.Try your best to be original, avoiding clichés or overused tropes. Do not use ornamental language and focus on nuance, simplicity, and subtext.Start directly with your response \n **{{instruction}}**

Table 14: Prompts for generating instructions and responses

A.1 Idiosyncracy Span Detection Prompt

Idiosyncracy Span Detection Prompt

You are given a paragraph of writing, and your goal is to provide feedback by selecting spans of text in the writing that could be improved and assign each problematic span to an error category. Below, we list the 7 error categories that you can choose from.
You are also provided 2 examples of paragraphs that were annotated by professional writers, which you can use to better understand the task and the error categories.
Error Categories:
- "Awkward Word Choice and Phrasing": Suggestions for better word choices or more precise phrasing to enhance clarity and readability.
- "Cliche": The use of hackneyed phrases or overly common imagery that lacks originality or depth.
- "Poor Sentence Structure": Feedback on the construction of sentences, recommending changes for better flow, clarity, or impact.
- "Unnecessary/Redundant Exposition": Redundant or non-essential parts of the text that could be removed/rephrased for conciseness.
- "Lack of Specificity and Detail": Need for more concrete details or specific information to enrich the text and make it more engaging.
- "Purple Prose": Identifying parts of the text that are seen as unnecessary ornamental and overly verbose.
- "Tense Consistency": Comments pointing out inconsistencies in verb tense that need to be addressed for uniformity.

Example 1: Input Text

Output:
Example Output in JSON format.

Example 2:
(Similar to example 1)
Rules:
- Number of Spans – You can provide feedback on multiple spans, and multiple spans can have the same category.
- Span must be verbatim – The span you select must be verbatim from the paragraph, otherwise, the feedback will not be provided to the user.
- No Overlap – Spans should not overlap, and one span should not include the other.
- Single Category – Each span should have exactly one category from the categories listed above.

Paragraph:

PARAGRAPH